

What Target-wise ABA Learning Does on Deterministic Causal Fixtures

Evidence from controlled probes H0–H7b

Working synthesis for discussion with Fabrizio

4 August 2026

Status and purpose. This is a supervisor-facing working document. It introduces the causal fixture, learner input, terminology, configurations, and experimental sequence before stating any cross-cutting findings. The controlled probes labelled H0–H7b are complete and analysed. The labels are internal shorthand only, and each probe is described in the text by its scientific question and design. No cross-cutting conclusions are asserted yet. The six Findings headings and the Implications heading at the end are intentionally empty placeholders for subsequent drafting and review.

Abstract

The investigation asks what target-wise ABA Learning does when given only a finite binary table sampled from a known causal data-generating process. The principal process is a binary collider $A \rightarrow C \leftarrow B$ with independent stochastic roots and a deterministic AND child $C = A \wedge B$. The graph, mechanisms, population distribution, finite sample, learner-visible encoding, learned ABA framework, and evaluator-only reference are kept separate throughout.

The experimental sequence is recorded under labels H0–H7b for reference. It begins with a baseline comparison of the published Greedy and non-deterministic configurations across all targets on one frozen table (H0). Later probes each introduce one controlled question: finite-sample support (H1), an irrelevant trailing variable (H2), background-knowledge order with a leading distractor (H3), brave versus cautious learning (H4–H5), folding-token budget and grounded sample size (H6), and the `relto` versus `sechk` assumption-introduction options (H7). This document records what each probe was intended to reveal, its principal signal, and the procedural explanation supported by the generated deltas and Prolog traces.

Contents

1	Purpose, scope, and organising question	3
1.1	Research question	3
1.2	Scope	3
1.3	Object of study	3
2	Conceptual framework and notation	4
2.1	Seven objects kept separate	4
2.2	ABA framework	4
2.3	ABA Learning task	5

2.4	Exact-value target-wise encoding	5
2.5	Brave and cautious learning	5
2.6	Greedy and non-deterministic folding	6
2.7	Assumption introduction: <code>relto</code> and <code>sechk</code>	6
3	Causal fixtures and evaluator ground truth	7
3.1	Principal fixture: binary AND collider (H0)	7
3.2	Evaluator-only mechanism reference	8
3.3	Frozen sample for the principal fixture	8
3.4	Controlled extensions: trailing and leading distractors (H2, H3)	9
4	Experimental design and comparison logic	10
4.1	Learner configurations	10
4.2	Strength of comparisons	10
4.3	Probe sequence	11
5	Experimental sequence: H0–H7b	12
5.1	H0: baseline on the principal AND collider	12
5.2	H1: nested prefix omitting the rare (0, 0, 0) atom	15
5.3	H2: trailing isolated variable under Greedy	16
5.4	H3: BK-leading isolated variable under nd	16
5.5	H4: brave versus cautious under nd	18
5.6	H5: Greedy under cautious versus brave learning	23
5.7	H6: explain cautious-nd search cost	23
5.8	H7: change assumption introduction from <code>relto</code> to <code>sechk</code>	26
6	Findings	28
6.1	Finding 1: Greedy solutions retain all represented BK predictors	28
6.2	Finding 2: nd solutions are organised around the first BK predictor	28
6.3	Finding 3: brave nd can solve predictor-ambiguous root targets	28
6.4	Finding 4: the folding-token ceiling repeats one search attempt	28
6.5	Finding 5: cautious-nd trace size grows with grounded sample size	30
6.6	Finding 6: <code>relto</code> and <code>sechk</code> produce different assumption-introduction paths . . .	30
7	Implications	30

1 Purpose, scope, and organising question

1.1 Research question

For a finite binary table generated by independent stochastic roots and deterministic non-root mechanisms, and for each observed variable taken in turn as the learning target: what ABA framework does ABALearn emit under the configurations studied here, and how does that framework relate to the evaluator-only mechanism reference and to the information present in the population and in the sample?

The purpose is to describe how the learner transforms the supplied table into rules, assumptions, and contraries, and to classify each emitted theory against a target-dependent evaluator reference. On the principal AND collider $A \rightarrow C \leftarrow B$, only the deterministic child C has a desired learned rule over its observed parents:

```
c(A) :- a_val_1(A), b_val_1(A).
```

The roots A and B are generated by independent stochastic mechanisms. There is no underlying deterministic rule relating either root to the other observed variables, and none should be learned.

Relative to that target-dependent reference, the outcomes distinguished in this document are:

- for a deterministic child: a rule that coincides with the mechanism;
- for a deterministic child: a sample-adequate but non-minimal rule;
- for a deterministic child: a defeasible representation that differs from the mechanism string;
- for a root or child: a rule that holds on the finite sample because a population counterexample is missing from that sample;
- for a root: an assumption construction that satisfies the selected learning semantics while leaving the target underdetermined by its predictors;
- search termination under the configured procedure, including timeout or `completed_no_solution`, which for a stochastic root target may be the intended outcome under the reference above.

These classes separate procedural exit status, finite-sample adequacy, and mechanism alignment under the target-dependent reference.

1.2 Scope

The evidence comes from three small binary fixtures, fixed seed-42 samples, and the learner configurations described below. Learning is target-wise: one frozen table is held fixed, and each observed variable is selected in turn as the learning target. Changing the target changes the positive and negative examples and the background-knowledge encoding of the remaining columns; the underlying rows remain the same sample.

Only the deterministic child has an evaluator-desired deterministic rule over its observed parents. Root-target runs are diagnostics of the same pipeline when the target itself is generated stochastically: they show how finite support, learning semantics, and search budget shape the emitted theory or search outcome when no underlying deterministic rule from the other variables should be learned.

1.3 Object of study

The pipeline under study is target-wise ABA Learning over encoded tabular data. For each target it receives background knowledge and labelled examples only, and it emits an ABA framework of rules, assumptions, and contraries under a selected learning semantics and search configuration.

The evidence recorded here therefore concerns learned rule structure, search behaviour, and the relation of those outputs to the evaluator-only mechanism reference for the fixture under study.

2 Conceptual framework and notation

This section fixes the objects, ABA notions, target-wise encoding, and learner controls used throughout the experimental sequence.

2.1 Seven objects kept separate

The experimental reasoning distinguishes the following objects:

1. a generating DAG G ;
2. structural mechanisms for the variables in G ;
3. the exact observational population distribution P^* ;
4. a finite IID sample $\mathcal{D}_n \sim P^*$;
5. target-specific learner input derived from $\mathcal{D}_n$;
6. the learned ABA framework and its emitted delta (the learned additions relative to the input background knowledge); and
7. evaluator-side judgements about sample adequacy, mechanism alignment, and minimality relative to the target-dependent reference.




Figure 1: Experimental separation. Only the finite table and its target-wise encoding enter the learner. The generating graph and mechanism reference are reserved for evaluation.

Ordinary faithfulness is a property of the pair (P^*, G) . Finite-sample adequacy is a property of a learned framework relative to $\mathcal{D}_n$ under the selected learning semantics. A parent literal appearing in a learned body is a syntactic fact about the emitted delta; any judgement that the delta aligns with the mechanism is an evaluator-side reading.

2.2 ABA framework

An ABA framework is a tuple

$$\mathcal{F} = \langle \mathcal{L}, \mathcal{R}, \mathcal{A}, \bar{\cdot} \rangle,$$

where $\mathcal{L}$ is a language, $\mathcal{R}$ is a set of rules, $\mathcal{A} \subseteq \mathcal{L}$ is a set of assumptions, and $\bar{\cdot}$ maps each assumption $\alpha \in \mathcal{A}$ to its contrary $\bar{\alpha} \in \mathcal{L}$.

A rule has the form

$$h \leftarrow b_1, \dots, b_m.$$

If $m = 0$, it is a fact. Write $\Delta \vdash_{\mathcal{R}} s$ when a claim $s \in \mathcal{L}$ is derivable from a set of assumptions $\Delta \subseteq \mathcal{A}$ using the rules in $\mathcal{R}$.

Attacks are generated by contraries: Δ attacks an assumption $\alpha \in \mathcal{A}$ when $\Delta \vdash_{\mathcal{R}} \bar{\alpha}$. A set $\Delta \subseteq \mathcal{A}$ is *conflict-free* if it attacks none of its own members. It is *closed* if whenever $\Delta \vdash_{\mathcal{R}} \alpha$ for some $\alpha \in \mathcal{A}$, then $\alpha \in \Delta$.

A set $\Delta \subseteq \mathcal{A}$ is a *stable extension* of $\mathcal{F}$ if it is closed, conflict-free, and attacks every assumption in $\mathcal{A} \setminus \Delta$. Write $\text{Stab}(\mathcal{F})$ for the collection of all stable extensions of $\mathcal{F}$.

A claim $s \in \mathcal{L}$ is accepted with respect to a stable extension Δ if there exists $\Delta' \subseteq \Delta$ such that $\Delta' \vdash_{\mathcal{R}} s$; write $\mathcal{F} \models_{\Delta} s$ when this holds. Write $\mathcal{F} \models_{\text{caut}} s$ when s is accepted with respect to every stable extension of $\mathcal{F}$. A learned framework may therefore combine ordinary rules with explicit exceptional conditions expressed by assumptions and contraries.

2.3 ABA Learning task

An ABA Learning task begins with background knowledge and two finite sets of examples:

$$E^+ = \{e_1^+, \dots, e_p^+\}, \quad E^- = \{e_1^-, \dots, e_q^-\}.$$

The learner applies symbolic transformations to construct a framework $\mathcal{F}'$ under which the examples satisfy the selected consequence condition. The learned object is itself an ABA framework.

The transformations relevant here are:

- **Rote learning.** Add example-specific rules from the positive examples.
- **Folding.** Replace example-specific conditions with reusable BK predicates, producing intensional rules.
- **Subsumption.** Remove a specialised rule when an appropriate more general rule already exists.
- **Assumption introduction.** Make an over-general rule defeasible by adding an assumption, then learn rules for its contrary to represent exceptions.

2.4 Exact-value target-wise encoding

In this investigation the learning task is instantiated from a tabular sample. For each row identifier i , every non-target variable is represented by exactly one value predicate. When c is the target, for example, the BK contains facts such as

```
a_val_1(9).
b_val_0(9).
```

and row 9 contributes either $c(9) \in E^+$ or $c(9) \in E^-$, according to its target value. For a binary target Y ,

$$E_Y^+ = \{Y(i) : Y_i = 1\}, \quad E_Y^- = \{Y(i) : Y_i = 0\}.$$

All non-target columns are available as BK predictors in their serialised order. The encoding interface is the same for every target. The evaluator-side desired object still depends on the target's generating role: only a deterministic child has a desired parent mechanism rule.

2.5 Brave and cautious learning

Brave and cautious learning are the two consequence conditions under which a learned framework $\mathcal{F}'$ may be required to cover E^+ and exclude E^- .

Brave learning requires an existential witness: there must exist some $\Delta \in \text{Stab}(\mathcal{F}')$ that simultaneously accepts every positive example and rejects every negative example,

$$\exists \Delta \in \text{Stab}(\mathcal{F}') \quad [\forall e \in E^+, \mathcal{F}' \models_{\Delta} e \wedge \forall e \in E^-, \mathcal{F}' \not\models_{\Delta} e].$$




Figure 2: Target-wise encoding. All targets use the same rows; the selected target determines the positive and negative examples and the remaining columns determine the BK.

Cautious learning requires every positive example to be a cautious consequence and every negative example to be excluded as a cautious consequence,

$$\forall e \in E^+, \mathcal{F}' \models_{\text{caut}} e, \quad \forall e \in E^-, \mathcal{F}' \not\models_{\text{caut}} e.$$

Thus the two modes quantify differently over $\text{Stab}(\mathcal{F}')$: brave learning uses one witnessing extension, while cautious learning uses acceptance in every stable extension.

2.6 Greedy and non-deterministic folding

Folding is controlled by `folding_mode`, together with companion options for selection and folding space. The two modes used here are as follows.

folding_mode(greedy). A rote rule is folded in one deterministic pass against matching background predicates, under `folding_selection(mgr)` and `folding_space(bk)`. This is the folding regime of the AAMAS configuration.

folding_mode(nd). Folding proceeds incrementally under a folding-token budget (`folding_steps`), with `folding_selection(any)` and `folding_space(all)`. This is the folding regime of the ECAI configuration and of the cautious nd configurations used later.

Each published configuration also sets learning semantics and `asm_intro`. When two named configurations are compared as wholes, several options change together.

2.7 Assumption introduction: `relto` and `sechk`

Assumption introduction is the transformation that makes an over-general folded rule defeasible by placing an assumption in its body and learning rules for the contrary of that assumption. In ABALearn it is controlled by the option `asm_intro`, which takes one of two values: `relto` or `sechk`. Both are used after a folded candidate fails the configured post-folding entailment check. They differ in how the first assumption-repair step is chosen.

asm_intro(relto). The learner first searches the current framework for an existing assumption that is relative to the body of the folded rule under consideration. If such an assumption is found, it is reused in the repaired rule. If none is found, the learner may generate a new assumption and continue along the subsequent generalisation path.

`asm_intro(sechk)`. The learner first generates a new provisional assumption for the folded rule and then subjects the resulting framework to the semantic-checking branch of generalisation (`gen5/gen6`), which inspects stable-extension consequences before deciding how the assumption is retained or revised.




Figure 3: The two `asm_intro` options as alternative first steps of assumption introduction.

3 Causal fixtures and evaluator ground truth

This section records the generating processes, population certificates, evaluator-only mechanism references, and frozen samples used by the probes below. Mathematical labels are A, B, C, D ; learner-safe identifiers are `a,b,c,d`. In emitted Prolog rules the capital argument (for example the `A` in `c(A)`) is the row identifier, not the causal variable A .

3.1 Principal fixture: binary AND collider (H0)

The principal fixture is the unshielded collider

$$G_0 : \quad A \longrightarrow C \longleftarrow B$$

with no A – B edge. The roots are mutually independent,

$$A \sim \text{Bernoulli}(4/5), \quad B \sim \text{Bernoulli}(7/10), \quad A \perp\!\!\!\perp B,$$

and the child is the deterministic AND mechanism

$$C = A \wedge B.$$

$$\Pr(A = 1) = 0.8 \quad \Pr(B = 1) = 0.7$$




Figure 4: Principal generating DAG (H0). Randomness enters through the independent roots; the non-root mechanism is deterministic.

The induced observational population has four positive-probability atoms:

A	B	C	$P^*(A, B, C)$
0	0	0	$3/50 = 0.06$
0	1	0	$7/50 = 0.14$
1	0	0	$6/25 = 0.24$
1	1	1	$14/25 = 0.56$

The remaining four binary assignments $(0, 0, 1)$, $(0, 1, 1)$, $(1, 0, 1)$, and $(1, 1, 0)$ are structural zeros.

Causal sufficiency is declared by fixture design: the root exogenous noise terms are mutually independent and there is no omitted common cause. The causal Markov condition holds by construction of the joint,

$$P^*(a, b, c) = P^*(a)P^*(b)P^*(c \mid a, b).$$

Ordinary faithfulness was verified exactly by auditing every singleton-pair conditional-independence query against d -separation. The sole population independence is $A \perp\!\!\!\perp B$; conditioning on the collider yields $A \not\perp\!\!\!\perp B \mid C$, and each parent is dependent on C .

Under ordinary conditional-independence semantics, the skeleton is $A - C - B$ and the unshielded collider $A \rightarrow C \leftarrow B$ is oriented in the CPDAG. The corresponding MEC contains only G_0 . This is the ordinary CI-based reading. Deterministic mechanisms can impose additional functional constraints beyond ordinary CI. The present document stays with the ordinary CI reading. The graph, certificates, and mechanism reference are evaluator-only and are withheld from ABALearn.

3.2 Evaluator-only mechanism reference

For the deterministic child C , the canonical positive-state exact-value rule is

```
c(A) :- a_val_1(A), b_val_1(A).
```

On this fixture the formal truth-table rule, the population-supported rule, and the rule supported by the frozen sample coincide on that string. That string is the evaluator reference against which learned theories for target c are compared. Comparison uses rule bodies, assumptions, contraries, and the target’s generating role together: a mechanism-aligned string is one form of evidence; an assumption-rich but sample-adequate theory is another object under the same reference.

The roots A and B are generated by independent stochastic mechanisms $A \sim \text{Bernoulli}(4/5)$ and $B \sim \text{Bernoulli}(7/10)$. Neither root is a deterministic function of the other observed variables, so when $\mathbf{a}$ or $\mathbf{b}$ is the learning target no underlying deterministic rule from those predictors should be learned.

3.3 Frozen sample for the principal fixture

The principal table is the target-free IID sample with $n = 30$ and seed 42:

Assignment	Count
$(1, 1, 1)$	17
$(1, 0, 0)$	7
$(0, 1, 0)$	4
$(0, 0, 0)$	2

All four population atoms appear. Under the binary exact-value encoding the implied example counts are

$$a : 24^+/6^-, \quad b : 21^+/9^-, \quad c : 17^+/13^-.$$

For child target **c**, every positive example has $(A, B) = (1, 1)$. For root target **a**, the predictor cell $(B, C) = (0, 0)$ contains both labels:

Row	A (target)	B	C	Class for target a
9	1	0	0	positive
26	0	0	0	negative

Rows 9 and 26 therefore share the same BK literals `b_val_0` and `c_val_0` while carrying opposite labels for A . No deterministic rule over (B, C) separates them. Target **b** has the symmetric ambiguity on predictor cells involving (A, C) .

3.4 Controlled extensions: trailing and leading distractors (H2, H3)




Figure 5: Controlled four-variable extensions. Dashed red nodes are independent isolated variables. Serialisation order induces the stated BK predictor order for target **c**.

Trailing isolated D (H2). This fixture preserves the H0 collider, the H0 root laws $A \sim \text{Bernoulli}(4/5)$ and $B \sim \text{Bernoulli}(7/10)$, and the mechanism $C = A \wedge B$. It adds an independent isolated root $D \sim \text{Bernoulli}(1/2)$ with no edges. Variable order places **d** last, so for target **c** the exact-value BK predictors appear as $[a, b, d]$: the distractor is trailing. The evaluator rule for **c** remains

```
c(A) :- a_val_1(A), b_val_1(A).
```

The variables A , B , and the isolated D are each generated stochastically. When any of them is the learning target, there is no underlying deterministic rule from the remaining observed variables that should be learned. The frozen table used below is again an IID sample with $n = 30$ and seed 42 on this four-variable process.

BK-leading isolated A (H3). This fixture uses independent roots

$$A \sim \text{Bernoulli}(1/2), \quad B \sim \text{Bernoulli}(4/5), \quad D \sim \text{Bernoulli}(7/10),$$

with A isolated (no edges) and deterministic child $C = B \wedge D$. Variable order places **a** first, so for target **c** the BK predictors appear as $[a, b, d]$: the distractor is leading. The evaluator rule for **c** is

```
c(A) :- b_val_1(A), d_val_1(A).
```

The isolated variable A and the roots B and D are each generated stochastically. When any of them is the learning target, there is no underlying deterministic rule from the remaining observed variables that should be learned. The frozen table used below is an IID sample with $n = 30$ and seed 42 on this process.

Both extensions have verified Markov factorisation and exact ordinary faithfulness certificates. As with the principal fixture, the generating graph, certificates, and mechanism references are withheld from the learner.

4 Experimental design and comparison logic

This section records the learner configurations used in the probes, how strongly each designed contrast supports attribution to a single change, and the ordered probe map. Detailed outcomes appear in the experimental sequence that follows.

4.1 Learner configurations

Identity	Semantics	Folding	Sel.	Space	Steps	asm_intro
aamas2025	brave	greedy	mgr	BK	—	relto
ecai2024	brave	nd	any	all	10	relto
baseline_cautious	cautious	nd	any	all	10	relto
greedy_cautious	cautious	greedy	mgr	BK	—	relto
nd_brave_sechk	brave	nd	any	all	10	sechk
nd_cautious_sechk	cautious	nd	any	all	10	sechk

Here “Steps” is `folding_steps`, which bounds the non-deterministic folding-token budget under nd folding. All six identities also enable the post-learning checked-ASP artefact (`check_ic`). The cautious nd identities additionally pin `post_folding_test_entailment(true)`, which is likewise the effective default under the brave configurations used here.

`aamas2025` and `ecai2024` are the repository presets corresponding to the AAMAS 2025 and ECAI 2024 learning-option bundles used in this work. `baseline_cautious`, `greedy_cautious`, `nd_brave_sechk`, and `nd_cautious_sechk` are experimental identities for controlled comparison. In particular, `greedy_cautious` matches `aamas2025` except for `learning_mode(cautious)`.

The H6a configurations match `baseline_cautious` except for `folding_steps`, with $M \in \{1, 2, 5, 10\}$. H7a matches `ecai2024` except `asm_intro(sechk)`. H7b matches `baseline_cautious` except `asm_intro(sechk)`.

4.2 Strength of comparisons

Each probe is read at one of three attribution strengths.

- **Exact paired option ablation.** Fixture, sample, encoding, and the remaining learner options are held fixed; one learning option changes. Probes H4 and H4b change `learning_mode` under otherwise fixed nd / `relto` settings; H5 changes `learning_mode` under the fixed Greedy / `mgr` / BK / `relto` bundle; H6a changes `folding_steps`; H7a and H7b change only `asm_intro`.

- **Controlled data or fixture intervention.** The learner configuration is held fixed (or compared only within one configuration); one designed property of the sample or fixture changes. H1 uses a nested seed-42 prefix $n = 25$ of the principal table under AAMAS; H2 adds a trailing isolated D under AAMAS; H3 uses the BK-leading isolated- A fixture under ECAI (with AAMAS as secondary); H6b varies nested sample size $n \in \{30, 60, 90\}$ at fixed `folding_steps(2)` under `baseline_cautious`.
- **Bundle-level contrast.** Several learner settings change together, so a behavioural difference is attributed to the configuration bundle rather than to one option. H0 compares `aamas2025` with `ecai2024` on the principal table. Where both exist, contrasts between `greedy_cautious` and `baseline_cautious` are likewise bundle-level (Greedy versus nd), as used secondarily in H5.

The strength level states how far a difference among outcomes, deltas, or traces can be attributed to the designed intervention of that probe.

4.3 Probe sequence

Probe	Question	Controlled intervention	What is inspected
H0	What do the Greedy and nd bundles learn for each target on the principal AND collider?	Principal fixture, frozen $n = 30$ seed 42; <code>aamas2025</code> and <code>ecai2024</code> ; targets <code>a,b,c</code>	Exit status, deltas, and traces by target and configuration.
H1	How does omitting the rare (0,0,0) atom affect AAMAS root learning?	Nested seed-42 prefix $n = 25$ versus $n = 30$; AAMAS; same stream	Root deltas and the presence or absence of opposite-label witness rows.
H2	What does Greedy do with an irrelevant trailing variable on the child?	Principal AND collider plus isolated D ; AAMAS; target <code>c</code>	Whether <code>d</code> appears in the learned bodies relative to the D -free evaluator rule.
H3	What does nd learning do when an isolated variable is first in BK?	BD AND collider with BK-leading isolated A ; order $[a, b, d]$; ECAI primary on target <code>c</code>	First fold, retained BK literals, and the final child theory.
H4	How does switching brave nd to cautious nd affect root targets on the principal table?	Learning mode only, under fixed nd / <code>relto</code> ; targets <code>a,b</code> with control <code>c</code>	Root exit status and the control child delta.
H4b	How does the same brave-to-cautious switch affect the child repair on the BK-leading fixture?	Learning mode only on H3 target <code>c</code> ; nd / <code>relto</code> fixed	Shared structure versus the contrary-rule close under <code>alpha_3</code> .
H5	How does brave-to-cautious behave inside the Greedy bundle?	<code>learning_mode</code> only under Greedy / <code>mgr</code> / BK / <code>relto</code> ; 18 paired cells across the completed AAMAS collections	Paired outcomes, deltas, traces, and runtime; secondary contrast to nd cautious where available.
H6a	How does the folding-token ceiling affect cautious-nd search cost on roots?	<code>folding_steps</code> $M \in \{1, 2, 5, 10\}$; principal table $n = 30$; <code>baseline_cautious</code>	Trace length, runtime, token bands, and control child behaviour.

Probe	Question	Controlled intervention	What is inspected
H6b	How does grounded sample size affect the same cautious-nd search?	Nested $n \in \{30, 60, 90\}$ at fixed $M = 2$; principal fixture; baseline_cautious	Search topology and per-band growth with n .
H7a	How do relto and sechk differ under brave nd?	asm_intro only; principal and BK-leading fixtures; seven target pairs	Paired deltas and the assumption-introduction control-flow fork.
H7b	How do relto and sechk differ under cautious nd?	asm_intro only; same two fixtures; baseline_cautious versus nd_cautious_sechk	Child deltas, root search cost, and whether runs complete within the wall-time limit.

5 Experimental sequence: H0–H7b

5.1 H0: baseline on the principal AND collider

Question and purpose. On the principal fixture $A \rightarrow C \leftarrow B$ with $C = A \wedge B$, frozen sample $n = 30$ seed 42, and exact-value encoding, what ABA frameworks do the Greedy bundle (aamas2025) and the nd bundle (ecai2024) emit when each of **a**, **b**, and **c** is the learning target?

Controlled comparison. The fixture, table, encoding, and target-specific examples were held fixed. The two configurations differ as full bundles: **aamas2025** uses brave learning with Greedy / **mgr** / BK folding and **relto**; **ecai2024** uses brave learning with nd / **any** / **all** folding and **relto**. The contrast is therefore bundle-level.

Configuration	Target	Outcome	Emphasised result
AAMAS	c	solved	Assumption-free rule c(A) :- a_val_1(A), b_val_1(A). , string-identical to the evaluator mechanism.
AAMAS	a,b	completed_no_solution	Both-zero co-occurrence candidates cover opposite-label rows. Contrary repair dead-ends under relto .
ECAI	c	solved	First fold uses BK-leading a_val_1 , repaired by an assumption attacked when b_val_0 .
ECAI	a,b	solved	Brave per-row assumption construction on the ambiguous predictor cell. A (resp. B) remains underdetermined by the other observed variables.

Child target **c**

Emphasised observation. Under AAMAS, every positive example for **c** has $(A, B) = (1, 1)$. Greedy folding with BK space therefore builds the single maximal co-occurrence body $\{\mathbf{a_val_1}, \mathbf{b_val_1}\}$. After mgr deduplication the emitted delta is

c(A) :- a_val_1(A), b_val_1(A).

Brave entailment succeeds immediately: no negative example shares that body, so no assumption is introduced. The string coincides with the evaluator mechanism reference for C .

Under ECAI, nd folding takes one BK predicate at a time. From an early positive, the first fold replaces the example identifier by `a_val_1`, the first predictor in BK order, yielding the unary rule `c(A) :- a_val_1(A)`. That body also holds of every negative with $A = 1, B = 0$. Assumption introduction produces

```
c(A) :- alpha_1(A), a_val_1(A).
c_alpha_1(A) :- b_val_0(A).
assumption(alpha_1(A)).
contrary(alpha_1(A), c_alpha_1(A)) :- assumption(alpha_1(A)).
```

Trace-supported explanation. The ECAI theory is sample-adequate under brave checking: c is derived when $a = 1$, unless the assumption is attacked when $b = 0$. Across the target rule and the contrary it cites both parents of C . Its ABA shape is nevertheless different from the AAMAS conjunction: the first fold used the leading BK predictor, and the second parent enters only through contrary learning.

Root targets a and b

Emphasised observation. For target **a**, positives fall into two predictor cells, $(B, C) = (1, 1)$ and $(B, C) = (0, 0)$. The cell $(B, C) = (0, 0)$ also contains negatives: row 9 is positive with joint values $(1, 0, 0)$, while rows 26 and 30 are negative with $(0, 0, 0)$, yet all share the BK literals `b_val_0` and `c_val_0`. Target **b** is the symmetric case with the roles of A and B swapped.

Under AAMAS, greedy folding proposes the two Horn candidates

```
a(A) :- b_val_1(A), c_val_1(A).
a(A) :- b_val_0(A), c_val_0(A).
```

The both-one rule is accepted. The both-zero rule covers the negative rows 26 and 30, so brave entailment fails. The learner introduces an assumption on that body. Contrary rote learning then cites those negatives, and greedy folding of the contrary reconstructs the same both-zero body. Under `asm_intro(relto)` a further relative assumption fails for that body. The engine reports KO, and the cell ends as `completed_no_solution` with no emitted delta. Target **b** mirrors this path.

Under ECAI, nd folding again proceeds one predicate at a time. The learner first builds a safe block for the both-one positives,

```
a(A) :- alpha_1(A), b_val_1(A).
c_alpha_1(A) :- c_val_0(A).
```

That block covers positives with $(B, C) = (1, 1)$ and excludes negatives with $B = 1$ and $C = 0$. The both-zero positives remain uncovered, so a second ordinary target rule is still required.

Folding an early both-zero positive (row 8) produces the candidate `a(A) :- b_val_0(A)`. Brave entailment fails on that candidate. The same BK literals also hold of the negative examples in rows 26 and 30. The learner therefore introduces an assumption α_2 on the body, giving `a(A) :- alpha_2(A), b_val_0(A)`. Contrary rote learning cites rows 26 and 30 as grounds for `c_alpha_2`.

Folding that contrary first tries `b_val_0`. Under `relto` this returns KO against the already present α_2 . The next fold tries `c_val_0`, yielding `c_alpha_2(A) :- c_val_0(A)`. Even that unary contrary remains too strong for the joint example set. The learner therefore introduces a further assumption α_3 on the contrary. Rote learning of `c_alpha_3` then cites the both-zero positives, including row 9. Folding that contrary re-uses α_2 together with `b_val_0`. The trace records this close as **found**:

alpha_2/1 ... OK. The final delta for target **a** is

```
a(A) :- alpha_1(A), b_val_1(A).
a(A) :- alpha_2(A), b_val_0(A).
c_alpha_1(A) :- c_val_0(A).
c_alpha_2(A) :- alpha_3(A), c_val_0(A).
c_alpha_3(A) :- alpha_2(A), b_val_0(A).
assumption(alpha_1(A)).
assumption(alpha_2(A)).
assumption(alpha_3(A)).
contrary(alpha_1(A),c_alpha_1(A)) :- assumption(alpha_1(A)).
contrary(alpha_2(A),c_alpha_2(A)) :- assumption(alpha_2(A)).
contrary(alpha_3(A),c_alpha_3(A)) :- assumption(alpha_3(A)).
```

Target **b** is the mirror obtained by swapping the roles of A and B .

Trace-supported explanation. The decisive geometry is the predictor cell $(B, C) = (0, 0)$. Row 9 is a positive example for target **a**. Its joint values are $(A, B, C) = (1, 0, 0)$. Row 26 is a negative example. Its joint values are $(0, 0, 0)$. Both rows carry the BK facts **b_val_0** and **c_val_0**. Any Horn rule that uses only those predictors therefore treats the two rows alike. It would derive **a** at both identifiers, or at neither.

The emitted α_2/α_3 fragment works because assumptions are grounded per row identifier. At a fixed row i , the relevant instances are $\alpha_2(i)$ and $\alpha_3(i)$, with contraries **c_alpha_2(i)** and **c_alpha_3(i)**. The local rules are

```
a(i) :- alpha_2(i), b_val_0(i).
c_alpha_2(i) :- alpha_3(i), c_val_0(i).
c_alpha_3(i) :- alpha_2(i), b_val_0(i).
```

On every both-zero row, both **b_val_0(i)** and **c_val_0(i)** hold. Accepting $\alpha_2(i)$ therefore derives **c_alpha_3(i)** and attacks $\alpha_3(i)$. Accepting $\alpha_3(i)$ derives **c_alpha_2(i)** and attacks $\alpha_2(i)$. The two assumptions cannot both be accepted at the same row.

This leaves two locally stable choices at each both-zero identifier. If $\alpha_2(i)$ is accepted and $\alpha_3(i)$ is rejected, then **a(i)** is derived. If $\alpha_3(i)$ is accepted and $\alpha_2(i)$ is rejected, then **a(i)** remains underived.

Brave learning asks for one witnessing stable extension Δ of the whole grounded framework. That extension must cover every positive example and exclude every negative example at once. Such a witness exists. On row 9 it accepts $\alpha_2(9)$ and rejects $\alpha_3(9)$. The rule **a(9) :- alpha_2(9), b_val_0(9)** then derives the positive example $(1, 0, 0)$. On row 26 it accepts $\alpha_3(26)$ and rejects $\alpha_2(26)$. The same BK facts are present. Without $\alpha_2(26)$ the rule for **a** leaves **a(26)** underived, so the negative example $(0, 0, 0)$ is excluded. The two rows are therefore separated by their ground assumption atoms, even though their BK predictors are identical.

Row	Joint (A, B, C)	Choice in Δ	Effect
9	$(1, 0, 0)$ (positive)	$\alpha_2(9)$ in, $\alpha_3(9)$ out	a(9) derived
26	$(0, 0, 0)$ (negative)	$\alpha_2(26)$ out, $\alpha_3(26)$ in	a(26) blocked

Other stable extensions may reverse those local choices. Brave entailment requires only that at least one joint witness exists. The framework therefore passes brave checking on the finite sample.

For this mechanism, target **a** should have no learned solution: A is generated stochastically, and there is no underlying deterministic rule from the other observed variables. A solution was nev-

ertheless found. Brave entailment permits one stable extension to assign opposite statuses to the ground assumptions $\alpha_2(i)$ and $\alpha_3(i)$ at different row identifiers. That per-row split covers the positive example $(1, 0, 0)$ and excludes the negative example $(0, 0, 0)$ even though the two rows share the same BK predictors. The ECAI **solved** outcome records that brave covering of the finite examples under the selected semantics.

Evidential boundary. H0 is the establishing bundle contrast on one complete-support table. Procedural **solved** covers both the AAMAS mechanism-string match for **c** and the ECAI assumption-rich theories for **c** and the roots. Attribution of the configuration difference is to the full AAMAS versus ECAI option bundle.

Primary evidence record: docs/experiments/qualitative/M13-C3-binary-collider-and/learning_analysis.md

5.2 H1: nested prefix omitting the rare $(0, 0, 0)$ atom

Question and purpose. Under the fixed AAMAS bundle on the principal AND collider, what happens to the root targets when the seed-42 table is replaced by its nested prefix of length 25, which omits the only sample atoms $(0, 0, 0)$?

Controlled comparison. Everything except sample size is held fixed. H0 uses the first 30 rows of the seed-42 stream. H1 uses the exact first 25 rows of that same stream. The omitted rows are CSV rows 26 and 30, the only H0 occurrences of $(0, 0, 0)$. The population still assigns mass $P^*(0, 0, 0) = 3/50$ to that atom.

Target	H0 ($n = 30$)	H1 ($n = 25$)
a,b	completed_no_solution	solved with both-one and both-zero Horn rules
c	solved with the evaluator AND string	solved with the same AND string

Emphasised observation. The H1 root delta for target **a** is

```
a(A) :- b_val_1(A), c_val_1(A).
a(A) :- b_val_0(A), c_val_0(A).
```

Target **b** emits the mirror pair. The both-zero rules are false in the population. They are sample-adequate on H1 only because the prefix contains no negative example in the predictor cell $(B, C) = (0, 0)$ (resp. $(A, C) = (0, 0)$).

Trace-supported explanation. H0 already showed that, under AAMAS, the both-zero body is proposed for root **a** and is rejected when rows 26 and 30 are present as negatives in that cell. H1 removes exactly those rows. The same body then faces only positives in that cell, so brave entailment accepts it and the run ends as **solved**.

For these stochastic roots, the intended reading is that no deterministic rule from the other observed variables should be learned. The H0 outcome **completed_no_solution** matches that reading. The H1 outcome **solved** is a spurious finite-sample covering created by omitting the rare blocking atom. Child **c** is unchanged and serves only as a control.

Evidential boundary. H1 is a nested-prefix support ablation under fixed AAMAS settings.

Primary evidence record: docs/experiments/qualitative/M13-C3-binary-collider-and/h1_support_ablation.md

5.3 H2: trailing isolated variable under Greedy

Question and purpose. On the principal AND collider extended by an independent isolated root D , does AAMAS learning of the deterministic child c retain irrelevant d literals, or does it return the evaluator D -free rule $c(A) :- a_val_1(A), b_val_1(A).?$

Controlled comparison. The H0 collider, root laws, and mechanism $C = A \wedge B$ are preserved. An isolated root $D \sim \text{Bernoulli}(1/2)$ is added with no edges. Variable order places d last, so for target c the BK predictors are $[a, b, d]$. The run uses AAMAS on the frozen $n = 30$ seed-42 sample. Among the eighteen positive c rows, both $D = 1$ and $D = 0$ occur nine times.

Emphasised observation. AAMAS solves target c and emits

```
c(A) :- a_val_1(A), b_val_1(A), d_val_1(A).  
c(A) :- a_val_1(A), b_val_1(A), d_val_0(A).
```

Irrelevant D is retained. The two rules are sample-adequate and jointly cover the observed values of D , but they are non-minimal relative to the evaluator mechanism

```
c(A) :- a_val_1(A), b_val_1(A).
```

On the three-variable H0 table without D , the same AAMAS bundle emitted exactly that D -free string.

Trace-supported explanation. Under greedy folding with `folding_space(bk)`, each positive example is folded into a maximal co-occurrence body from the BK literals that hold of that row. A positive with $D = 1$ therefore keeps `d_val_1`. A positive with $D = 0$ keeps `d_val_0`. The trace shows this explicitly: after `a_val_1` and `b_val_1`, the next fold attaches the matching `d_val_*` literal. Mgr selection then keeps the resulting bodies. There is no later step that merges the complementary pair by deleting the d literal. The failure mode under test is therefore distractor retention on a solved child, not failure to solve c .

Evidential boundary. H2 is a controlled fixture intervention under fixed AAMAS settings. The primary signal is the balanced $n = 30$ distractor split on target c .

Primary evidence record: docs/experiments/qualitative/M13-C3-binary-collider-and/h2_irrelevant_covariate.md

5.4 H3: BK-leading isolated variable under nd

Question and purpose. On a BD AND collider with an isolated root placed first in BK order, does ECAI learning of the deterministic child c select that isolated variable in the first fold and retain it in the final theory, or does it return the evaluator rule $c(A) :- b_val_1(A), d_val_1(A).?$

Controlled comparison. The fixture is $B \rightarrow C \leftarrow D$ with $C = B \wedge D$ and an independent isolated root $A \sim \text{Bernoulli}(1/2)$. Variable order places **a** first, so for target **c** the BK predictors are $[a, b, d]$. The primary run is ECAI brave nd on the frozen $n = 30$ seed-42 sample. An AAMAS run on the same table supplies a bundle-level contrast. Among the fifteen positive **c** rows, every row has $(B, D) = (1, 1)$, while both values of A occur.

Emphasised observation. The ECAI trace opens with the first fold

```
folding result: c(A) <- [a_val_1(A)]
```

Assumption introduction follows immediately. The full emitted delta for target **c** is:

```
c(A) :- alpha_1(A), a_val_1(A).
c(A) :- alpha_2(A), a_val_0(A).
c_alpha_1(A) :- b_val_0(A).
c_alpha_1(A) :- alpha_3(A), b_val_1(A).
c_alpha_2(A) :- d_val_0(A).
c_alpha_2(A) :- b_val_0(A).
c_alpha_3(A) :- alpha_1(A), a_val_1(A).
assumption(alpha_1(A)).
assumption(alpha_2(A)).
assumption(alpha_3(A)).
contrary(alpha_1(A), c_alpha_1(A)) :- assumption(alpha_1(A)).
contrary(alpha_2(A), c_alpha_2(A)) :- assumption(alpha_2(A)).
contrary(alpha_3(A), c_alpha_3(A)) :- assumption(alpha_3(A)).
```

The two ordinary target rules are conditioned only on value literals of the isolated variable A (together with assumptions α_1 and α_2). Mechanism parents enter through contrary rules: **c_alpha_1** and **c_alpha_2** cite b and d , and under the $a=1$ branch one contrary rule of α_1 contains a further assumption α_3 . The theory is therefore organised around A , with B and D used to repair over-general A -conditioned target rules. The evaluator reference for this target remains **c(A) :- b_val_1(A), d_val_1(A) ..**

On the same sample, AAMAS also solves **c** while retaining **a**:

```
c(A) :- a_val_1(A), b_val_1(A), d_val_1(A).
c(A) :- a_val_0(A), b_val_1(A), d_val_1(A).
```

Trace-supported explanation. Under nd folding with **any** selection, the first admissible BK predicate becomes the first generalising fold. Here that predicate is an **a**-value literal, because **a** leads the serialised BK. The unary fold over-generalises, so the learner introduces α_1 (then α_2) and learns contraries for those A -conditioned target rules. Later folds place b and d in **c_alpha_1** and **c_alpha_2**, but the ordinary target rules keep the initial a -value literals. The AAMAS contrast reaches distractor retention by a different route: greedy maximal co-occurrence keeps every matching literal, including leading **a**, and splits the AND across the two values of A .

In both cases the child is solved, and in both cases the isolated leading variable is retained. The failure mode is distractor inclusion driven by BK order.

Contrary-structure divergence under the two A -conditioned target rules. The ECAI delta is asymmetric under the two ordinary target rules conditioned on $a=1$ and $a=0$. That asymmetry is part of the H3 theory and deserves a separate reading.

Under $a=1$ (α_1), the contrary of α_1 is *assumption-bearing*: one contrary rule contains a further

assumption α_3 :

```
c_alpha_1(A) :- b_val_0(A).  
c_alpha_1(A) :- alpha_3(A), b_val_1(A).  
c_alpha_3(A) :- alpha_1(A), a_val_1(A).
```

Under $a=0$ (α_2), the contrary of α_2 is *ordinary*: every contrary rule uses only BK literals, citing the mechanism parents directly:

```
c_alpha_2(A) :- d_val_0(A).  
c_alpha_2(A) :- b_val_0(A).
```

So the $a=1$ side is repaired by a further assumption α_3 , while the $a=0$ side is repaired by ordinary b/d literals alone.

The divergence is a search-order effect under nd folding with **any** selection, *b-before-d* BK order after leading **a**, cover-and-repair, and **asm_intro(relto)**. The underlying mechanism of *C* is the same on both sides of the split. The same assumption-bearing α_1 contrary and ordinary α_2 contrary appear under both H3 brave and H4b cautious. Concretely:

1. Contrary repair for **c_alpha_1** occurs first, because the first positive example has $a=1$.
2. While repairing that contrary, the first contrary seed has $(b, d) = (0, 1)$, so search commits to **b_val_0**. The remaining contrary residuals have $(b, d) = (1, 0)$.
3. Folding those residuals with **b_val_1** over-generalises. Cover-and-repair introduces α_3 on the **b_val_1** body.
4. When search later repairs **c_alpha_2** under $a=0$, a fold that would reuse **b_val_1** meets the already declared α_3 under **relto** and is rejected. Search then falls through to **d_val_0** and **b_val_0**, yielding the ordinary BK-only contrary.

Thus the assumption-bearing α_1 contrary and the ordinary α_2 contrary are two stages of one sequential repair path. The $a=0$ side inherits the assumption minted while repairing the $a=1$ side.

Evidential boundary. H3 is a controlled fixture intervention. The primary signal is ECAI first-fold retention of BK-leading isolated *A* on target **c**. The assumption-bearing versus ordinary contrary split under α_1 and α_2 is a secondary procedural feature of that same theory.

Primary evidence record: docs/experiments/qualitative/M13-C3-binary-collider-and/h3_bk_leading_distractor.md

5.5 H4: brave versus cautious under nd

5.5.1 H4: roots on the H0 table

Question and purpose. On the principal AND-collider table of H0, does changing **learning_mode(brave)** to **learning_mode(cautious)** alter the outcomes for root targets **a** and **b** that, under **ecai2024**, relied on a residual per-row assumption construction separating rows 9 and 26, and what happens to control target **c**?

Controlled comparison. The H0 fixture, frozen $n = 30$ seed-42 table, target-specific examples, and BK were held fixed. The two learner configurations compared are **ecai2024** and **baseline_cautious**. Both use nd folding with **folding_selection(any)**, **folding_space(all)**, **asm_intro(relto)**, and **folding_steps(10)**. They differ in learning mode: **ecai2024** sets **learning_mode(brave)**; **baseline_cautious** sets **learning_mode(cautious)**. The locked **ecai2024** outputs are the brave comparator.

Target	ecai2024	baseline_cautious
<i>a</i>	solved	completed_no_solution
<i>b</i>	solved	completed_no_solution
<i>c</i>	solved	solved, identical delta

Emphasised observation. For target *a*, the `ecai2024` and `baseline_cautious` traces coincide through the early $\alpha_1, \alpha_2, \alpha_3$ construction from H_0 . They diverge at one assumption-introduction step. After

```
folding result: c_alpha_3(A) <- [b_val_0(A)]
```

the learner proposes to re-use the already declared assumption α_2 in that body:

```
c_alpha_3(A) :- alpha_2(A), b_val_0(A).
```

Under `ecai2024` the proposal is accepted; under `baseline_cautious` it is rejected:

```
% ecai2024:
found: alpha_2/1 ... OK, assumption introduction result:
  c_alpha_3(A) <- [alpha_2(A), b_val_0(A)]

% baseline_cautious:
found: alpha_2/1 ... KO, cannot introduce an assumption for [b_val_0(A)]
```

Under `ecai2024`, that OK completes the both-zero assumption-contrary fragment from H_0 and the learner writes a solution. Under `baseline_cautious`, the KO is followed by further failed folds, exhaustion of `folding_steps(10)`, and `completed_no_solution` with no emitted delta. Target *b* exhibits the same divergence with the roles of *A* and *B* swapped.

Under `baseline_cautious`, control target *c* remains solved. Its delta is identical to the locked `ecai2024` delta for *c*:

```
c(A) :- alpha_1(A), a_val_1(A).
c_alpha_1(A) :- b_val_0(A).
assumption(alpha_1(A)).
contrary(alpha_1(A), c_alpha_1(A)) :- assumption(alpha_1(A)).
```

Thus `learning_mode(cautious)` still admits an assumption-bearing theory for the child on this table, while rejecting the residual α_2 reuse for the roots.

Trace-supported explanation. Under both configurations the search builds toward the same both-zero assumption-contrary fragment:

```
a(A) :- alpha_2(A), b_val_0(A).
c_alpha_2(A) :- alpha_3(A), c_val_0(A).
c_alpha_3(A) :- alpha_2(A), b_val_0(A).
```

Write $\mathcal{F}_{bz}$ for a framework containing this fragment. The decisive conflict is observational. Row 9 contributes $a(9) \in E^+$ with joint values $(A, B, C) = (1, 0, 0)$. Row 26 contributes $a(26) \in E^-$ with joint values $(0, 0, 0)$. Relative to BK the two rows are indistinguishable: both carry `b_val_0` and `c_val_0`. No Horn body over those predictors alone can accept one row and reject the other.

The fragment separates the rows by ground assumption membership rather than by BK. At any

both-zero row identifier i , the assumptions $\alpha_2(i)$ and $\alpha_3(i)$ attack each other through their contrary rules, so a stable extension can take only one of the two local choices:

Choice at row i	Acceptance of $a(i)$	Derived contrary
$\alpha_2(i) \in \Delta, \alpha_3(i) \notin \Delta$	$\mathcal{F}_{bz} \models_{\Delta} a(i)$	<code>c_alpha_3(i)</code>
$\alpha_2(i) \notin \Delta, \alpha_3(i) \in \Delta$	$\mathcal{F}_{bz} \not\models_{\Delta} a(i)$	<code>c_alpha_2(i)</code>

Brave entailment. Under `learning_mode(brave)`, a learned framework $\mathcal{F}'$ is accepted when there exists one witnessing stable extension covering the examples:

$$\exists \Delta \in \text{Stab}(\mathcal{F}') [\forall e \in E^+, \mathcal{F}' \models_{\Delta} e \wedge \forall e \in E^-, \mathcal{F}' \not\models_{\Delta} e].$$

For $\mathcal{F}_{bz}$ such a witness exists. On row 9 it chooses $\alpha_2(9)$ in and $\alpha_3(9)$ out, so

```
a(9) :- alpha_2(9), b_val_0(9)
```

yields $\mathcal{F}_{bz} \models_{\Delta} a(9)$. On row 26 it chooses $\alpha_2(26)$ out and $\alpha_3(26)$ in, so $\mathcal{F}_{bz} \not\models_{\Delta} a(26)$. The closing rule

```
c_alpha_3(A) :- alpha_2(A), b_val_0(A).
```

is part of making those local choices coherent: whenever $\alpha_2(i)$ is accepted, `c_alpha_3(i)` is derived and attacks $\alpha_3(i)$.

Thus $\mathcal{F}_{bz}$ passes brave entailment because coverage may depend on different ground assumption choices at different row identifiers inside a single witness. The assumption-introduction check that proposes re-using α_2 on `c_alpha_3(A) <- [b_val_0(A)]` therefore returns OK, and `ecai2024` finishes with `solved`.

Failure under cautious entailment. Under `learning_mode(cautious)`, the same framework must satisfy the stronger condition

$$\forall e \in E^+, \mathcal{F}' \models_{\text{caut}} e, \quad \forall e \in E^-, \mathcal{F}' \not\models_{\text{caut}} e,$$

where $\mathcal{F}' \models_{\text{caut}} s$ means that s is accepted with respect to every stable extension of $\mathcal{F}'$.

Apply that test to $\mathcal{F}_{bz}$ at row 9. The choice that covers the positive is $\alpha_2(9) \in \Delta$. The opposite choice, $\alpha_3(9) \in \Delta$ and $\alpha_2(9) \notin \Delta$, also belongs to $\text{Stab}(\mathcal{F}_{bz})$. In any extension Δ' that takes the opposite choice, $\mathcal{F}_{bz} \not\models_{\Delta'} a(9)$. Therefore $a(9)$ fails $\mathcal{F}_{bz} \models_{\text{caut}} a(9)$. The brave cover of the positive exists as an existential cover. Cautious entailment requires universal cover, and $\mathcal{F}_{bz}$ leaves $a(9)$ undecided across stable extensions on this underdetermined BK cell.

That is why, under `baseline_cautious`, the same proposal returns

```
found: alpha_2/1 ... K0, cannot introduce an assumption for [b_val_0(A)]
```

The learner rejects the clause that would close the brave construction. Search then exhausts `folding_steps(10)` and returns `completed_no_solution`.

In short, `ecai2024` succeeds because brave entailment needs only some stable extension that assigns opposite α_2/α_3 statuses to rows 9 and 26. Under `baseline_cautious`, $a(9)$ fails to hold in every stable extension of $\mathcal{F}_{bz}$, so the residual α_2 reuse fails the cautious entailment check.

The control on target `c` is consistent with these semantics. Its contrary folds to the ordinary body `b_val_0`. That body treats rows that share the same BK literals uniformly, so the cautious entailment checks succeed, and the emitted delta matches the `ecai2024` delta for `c`.

Evidential boundary. H4 compares learner configurations `ecai2024` and `baseline_cautious` on the locked H0 table, with matched `nd` folding options, matched `asm_intro(relto)`, and `folding_steps(10)`. Under `baseline_cautious` the roots end as `completed_no_solution`.

Primary evidence record: `docs/experiments/qualitative/M13-C3-binary-collider-and/h4_cautious_vs_brave.md`

5.5.2 H4b: child c on the H3 table

Question and purpose. On the H3 BD AND table with BK-leading isolated A , under matched `nd` and `relto` settings, how does changing `learning_mode(brave)` to `learning_mode(cautious)` alter the contrary of α_3 when learning target c ?

Controlled comparison. The H3 fixture `m13_bucket3_binary_bd_and_lead_a`, frozen $n = 30$ seed-42 sample, target c , and BK order $[a, b, d]$ were held fixed. The two learner configurations compared are `ecai2024` and `baseline_cautious`. Both use `nd` folding with `folding_selection(any)`, `folding_space(all)`, `asm_intro(relto)`, and `folding_steps(10)`. They differ in learning mode: `ecai2024` sets `learning_mode(brave)`; `baseline_cautious` sets `learning_mode(cautious)`. The locked H3 `ecai2024` cell for target c is the brave comparator.

Object	<code>ecai2024</code>	<code>baseline_cautious</code>
Outcome for c	<code>solved</code>	<code>solved</code>
Ordinary target rules	$\alpha_1, a=1$ / $\alpha_2, a=0$	<code>same</code>
<code>c_alpha_1</code>	$b=0$; also $\alpha_3, b=1$	<code>same</code>
<code>c_alpha_2</code>	$d=0$; $b=0$	<code>same</code>
<code>c_alpha_3</code>	$\alpha_1, a=1$	$d=1$

Emphasised observation. Both configurations solve target c . Through the introduction of α_3 , their theories coincide with the H3 theory already recorded: ordinary target rules conditioned on leading a , an assumption-bearing contrary under $a=1$, and an ordinary b/d contrary under $a=0$. Exactly one closing rule differs. Under `ecai2024`:

```
c_alpha_3(A) :- alpha_1(A), a_val_1(A).
```

Under `baseline_cautious`:

```
c_alpha_3(A) :- d_val_1(A).
```

The traces diverge at the assumption-introduction check on that close. After

```
folding result: c_alpha_3(A) <- [a_val_1(A)]
```

`ecai2024` re-uses α_1 and stops:

```
found: alpha_1/1 ... OK, assumption introduction result:
c_alpha_3(A) <- [alpha_1(A), a_val_1(A)]
```

`baseline_cautious` rejects that reuse, rejects a subsequent attempt to re-use α_3 on `b_val_1`, and then accepts an ordinary `d_val_1` fold:

```
found: alpha_1/1 ... K0, cannot introduce an assumption for [a_val_1(A)]
```

```
folding result: c_alpha_3(A) <- [b_val_1(A)]
found: alpha_3/1 ... KO, cannot introduce an assumption for [b_val_1(A)]
folding result: c_alpha_3(A) <- [d_val_1(A)]
```

The full `baseline_cautious` delta is therefore the shared H3 theory with that single close replaced:

```
c(A) :- alpha_1(A), a_val_1(A).
c(A) :- alpha_2(A), a_val_0(A).
c_alpha_1(A) :- b_val_0(A).
c_alpha_1(A) :- alpha_3(A), b_val_1(A).
c_alpha_2(A) :- d_val_0(A).
c_alpha_2(A) :- b_val_0(A).
c_alpha_3(A) :- d_val_1(A).
assumption(alpha_1(A)).
assumption(alpha_2(A)).
assumption(alpha_3(A)).
contrary(alpha_1(A),c_alpha_1(A)) :- assumption(alpha_1(A)).
contrary(alpha_2(A),c_alpha_2(A)) :- assumption(alpha_2(A)).
contrary(alpha_3(A),c_alpha_3(A)) :- assumption(alpha_3(A)).
```

Trace-supported explanation. The shared prefix is the H3 theory under leading a . As recorded in H3, that prefix is produced by nd folding with **any** selection, b -before- d BK order, cover-and-repair, and `asm_intro(relto)`. It is reproduced under both learning modes through α_3 introduction. The assumption-bearing versus ordinary contrary split under α_1 and α_2 is therefore already present before the learning-mode contrast takes effect.

The learning-mode contrast is confined to the attack on α_3 . Under `ecai2024`, the learner closes that attack by residual reuse of α_1 :

```
c_alpha_3(A) :- alpha_1(A), a_val_1(A).
```

This clause couples the contrary of α_3 to the same assumption that appears in the ordinary target rule under $a=1$. Brave entailment asks only for some witnessing stable extension of the candidate framework that covers E^+ and excludes E^- . That residual reuse passes the brave assumption-introduction check, so `ecai2024` accepts it and finishes without folding a d -literal on `c_alpha_3`.

Under `baseline_cautious`, the entailment condition is universal: every positive must be a cautious consequence, and every negative must be excluded as a cautious consequence. The residual reuse of α_1 on `c_alpha_3` fails that test, exactly as the residual reuse of α_2 failed under H4: acceptance of the proposed extension depends on assumption choices that fail to hold in every stable extension. The introduction check therefore returns KO. On this H3 child cell an ordinary BK alternative remains available. Folding the contrary examples of α_3 with `d_val_1` satisfies the cautious entailment condition on $\langle E^+, E^- \rangle$, so the learner accepts

```
c_alpha_3(A) :- d_val_1(A).
```

and finishes with `solved`.

Relative to H4, the same learning-mode change therefore has the same local effect of rejecting a residual assumption reuse, together with a different global outcome. On the H0 roots, search exhausts `folding_steps(10)` without an accepted covering theory. On this H3 child cell, an ordinary parent literal closes the contrary of α_3 , while the ordinary target rules conditioned on leading a and the asymmetric contrary structure from H3 remain in place.

Evidential boundary. H4b compares `ecai2024` and `baseline_cautious` on the locked H3 target-c cell, with matched nd folding options and `asm_intro(relto)`. The emitted theories retain BK-leading isolated A in the ordinary target rules. The learning-mode effect isolated here is the body of `c_alpha_3`.

Primary evidence record: `docs/experiments/qualitative/M13-C3-binary-collider-and/h4b_cautious_split_under_a.md`

5.6 H5: Greedy under cautious versus brave learning

Question and purpose. Under otherwise fixed Greedy options, does changing `learning_mode(brave)` to `learning_mode(cautious)` alter the accepted or rejected theories on the completed Greedy collections from H0–H3?

Controlled comparison. The locked `aamas2025` collections for H0 ($n = 30$), H1 ($n = 25$), H2 ($n = 30$), the extreme H2 edge ($n = 2$), and H3 ($n = 30$) were re-run under the experimental configuration `greedy_cautious`. That configuration matches `aamas2025` on Greedy folding, `folding_selection(mgr)`, `folding_space(bk)`, `asm_intro(relto)`, and `check_ic`, and differs only by setting `learning_mode(cautious)`. This gives eighteen target-wise pairs on shared tables, BK, and examples. Configuration `greedy_cautious` is an experimental preset for this contrast.

Emphasised observation. On all eighteen pairs, `greedy_cautious` matches `aamas2025` on every outcome, and every solved delta is identical. The only systematic difference is runtime: cautious runs take approximately 1.6–2.3 times as long as the corresponding brave runs (mean ratio about 1.84), while remaining below one second.

Trace-supported explanation. Accepted and rejected Greedy theories are identical across the learning-mode change on these cells. The runtime increase comes from the entailment checks performed along the shared search path. Brave checks ask whether some answer set satisfies the joint example constraints. Cautious checks ask Clingo for cautious consequences (`-enum-mode=cautious`), that is atoms true in every answer set, and then test E^+ and E^- against that intersection. The same sequence of checks is therefore more expensive under cautious learning.

Evidential boundary. H5 is restricted to these eighteen Greedy pairs under matched non-mode options.

Primary evidence record: `docs/experiments/qualitative/M13-C3-binary-collider-and/h5_greedy_cautious.md`

5.7 H6: explain cautious-nd search cost

5.7.1 H6a: folding-token ceiling

Question and purpose. On the H0 table under `baseline_cautious`, how do root no-solution trace length and runtime scale with the folding-token ceiling $M = \text{folding_steps}$?

Controlled comparison. The H0 fixture, frozen $n = 30$ seed-42 sample, target-specific examples, BK, `learning_mode(cautious)`, nd folding, `folding_selection(any)`, `folding_space(all)`, and `asm_intro(relto)` were held fixed. The single varied option was `folding_steps(M)` for

$M \in \{1, 2, 5, 10\}$. Configurations `baseline_cautious_steps1`, `baseline_cautious_steps2`, and `baseline_cautious_steps5` match `baseline_cautious` except for M . The $M = 10$ collection reuses the locked H4 run under `baseline_cautious`. Roots `a` and `b` are the intended no-solution targets (stochastic; no desired deterministic rule from the other observed variables). Child `c` is the control.

M	Target	Outcome	Runtime (s)	Trace lines	N_{asm}
1	<code>a</code>	<code>completed_no_solution</code>	6.638	549	9
2	<code>a</code>	<code>completed_no_solution</code>	13.117	1047	18
5	<code>a</code>	<code>completed_no_solution</code>	32.775	2541	45
10	<code>a</code>	<code>completed_no_solution</code>	69.952	5031	90

Here N_{asm} is the number of stdout lines matching `generating NEW assumption: ...` in one target-wise Prolog trace (the literal engine message; “NEW” is part of that string). It counts assumption-*introduction events*, not the number of assumptions that remain in an accepted delta.

Target `b` mirrors these figures. Control `c` solves at every M inside `tokens(1)`: 134-line trace, about 1.3s, and a identical delta across all four ceilings.

Emphasised observation. For roots `a` and `b` the outcome is `completed_no_solution` at every tested M : no covering theory is accepted. Each run is a stack of M failed bands. The $M = 1$ trace has 549 lines and $N_{\text{asm}} = 9$. Each extra allowance adds 498 lines and nine further such events. At $M = 10$ the trace prints nine restarts

```
* Increasing folding tokens to: 2
...
* Increasing folding tokens to: 10
```

and ends with `* No solution found!`. The counts fit

$$\text{lines}(M) = 549 + 498(M - 1) = 51 + 498M, \quad N_{\text{asm}}(M) = 9M.$$

Runtime grows in the same way. Raising M multiplies failed bands. Every tested ceiling ends as `completed_no_solution`. The long H4 root traces are the $M = 10$ endpoint of this ladder.

Trace-supported explanation. A failed band is one complete cautious search at fixed token allowance T that exhausts without writing a delta. On these roots the H4 residual assumption reuse already returns the engine message `K0, cannot introduce an assumption for ...` inside the first allowance, so that first band fails.

Under `nd` folding, `folding_steps(M)` is the ceiling on how many such attempts may run. Search starts at `tokens(1)`. When the attempt at allowance T fails completely and $T + 1 \leq M$, the controller prints `* Increasing folding tokens to: \{T+1\}` and restarts the same learning problem on the same rules and the same E^+, E^- . The target still admits no cautiously acceptable cover, so each restart fails in the same way. M allowances therefore contribute M copies of the same failed band, which is why trace length and N_{asm} are linear in M .

Control `c` solves under `tokens(1)` and is therefore invariant in M .

Evidential boundary. H6a is a single-option ablation of `folding_steps` under `baseline_cautious` on the locked H0 $n = 30$ table, for these intended no-solution roots over $M \in \{1, 2, 5, 10\}$.

5.7.2 H6b: nested sample size

Question and purpose. On the H0 table under `baseline_cautious` with folding-token ceiling fixed at $M = 2$, how do root no-solution trace length and runtime scale when the frozen seed-42 sample is enlarged by nested prefixes $n \in \{30, 60, 90\}$?

Controlled comparison. The H0 fixture, seed 42, target-specific encoding, and learner configuration `baseline_cautious_steps2` were held fixed: `learning_mode(cautious)`, `nd folding`, `folding_selection(any)`, `folding_space(all)`, `asm_intro(relto)`, and `folding_steps(2)`. The single varied factor was sample size: nested prefixes $n = 30 \subset n = 60 \subset n = 90$ of the same seed-42 draw. Roots **a** and **b** are the intended no-solution targets; child **c** is the control. The $n = 30$ cell reuses the H6a $M = 2$ collection.

n	Target	Outcome	Runtime (s)	Trace lines	N_{asm}	N_{KO}
30	<i>a</i>	<code>completed_no_solution</code>	13.117	1047	18	32
60	<i>a</i>	<code>completed_no_solution</code>	22.169	1641	18	32
90	<i>a</i>	<code>completed_no_solution</code>	32.291	2241	18	32

Here N_{asm} is as in H6a: the count of stdout lines `generating NEW assumption: ....` N_{KO} is the count of stdout lines `KO, cannot introduce an assumption for ....` Both are event counts from one target-wise Prolog trace. They are included as a check that the *sequence of assumption-introduction and failed residual reuse events* is the same at every n , while trace length and runtime may still grow.

Target *b* mirrors these figures. Every root run prints a single restart to `tokens(2)`. Control **c** solves at every n with a identical delta; its trace grows only from 134 to 200 to 266 lines (about 1.3–3.0 s).

All four H0 population atoms are already present at $n = 30$. Their multiplicities, ordered as $(0, 0, 0)$, $(0, 1, 0)$, $(1, 0, 0)$, $(1, 1, 1)$, are

$$\begin{aligned}
 n = 30 &: (2, 4, 7, 17), \\
 n = 60 &: (4, 10, 12, 34), \\
 n = 90 &: (5, 17, 18, 50).
 \end{aligned}$$

Emphasised observation. Root outcomes are invariant in n : always `completed_no_solution`. The assumption-event counts are likewise invariant: $N_{\text{asm}} = 18$ and $N_{\text{KO}} = 32$ at every tested n , with one increase to `tokens(2)`. That matches H6a at $M = 2$, where each failed band contributes nine `generating NEW assumption` events, so two bands contribute $N_{\text{asm}} = 9M = 18$. What grows is how much text each band prints: target-*a* trace length rises by 594 then 600 lines for each additional 30 rows, and runtime rises somewhat faster than lines. Control **c** keeps the same accepted theory while its short trace lengthens mildly with n .

Trace-supported explanation. H6a varies how many failed bands may run. H6b holds that band count fixed at two allowances and varies how much grounded work sits inside each band. The fixed N_{asm} and N_{KO} show that the learner still takes the same pattern of assumption-introduction events and failed residual reuses; only the volume of grounding work between those events changes. Additional IID rows add example identifiers, rote facts, entailment and subsumption checks, and contrary batches. Support presence is already complete at $n = 30$, and only multiplicities change. The near-linear growth in trace length is therefore per-row grounding and bookkeeping cost. Runtime grows a little faster because ASP consequence checks scale with that larger grounding.

Evidential boundary. H6b is a nested- n ablation at fixed $M = 2$ under `baseline_cautious` on the locked H0 seed-42 prefixes $n \in \{30, 60, 90\}$. It measures cost when multiplicities change while assumption-event counts N_{asm} and N_{KO} stay fixed. The $n \times M$ grid was left for later work.

Primary evidence record: `docs/experiments/qualitative/M13-C3-binary-collider-and/h6_folding_and_n_ablation.md`

5.8 H7: change assumption introduction from `relto` to `sechk`

H7 holds the nd folding bundle fixed and changes only `asm_intro`. H7a uses brave learning. H7b uses cautious learning. The two probes therefore isolate the same control-flow fork under the two entailment modes already contrasted in H4/H4b.

5.8.1 H7a: brave nd

Question and purpose. Under brave nd on the locked H0 and H3 $n = 30$ tables, if only `asm_intro` changes from `relto` to `sechk`, which BK predictors appear in the accepted deltas?

Controlled comparison. Fixture, frozen seed-42 sample, target, examples, BK, `learning_mode(brave)`, nd folding, `folding_selection(any)`, `folding_space(all)`, `folding_steps(10)`, and `check_ic` were held fixed. Locked `ecai2024` uses `asm_intro(relto)`. Experimental `nd_brave_sechk` matches it except for `asm_intro(sechk)`. Seven pairs: H0 `a,b,c` and H3 `a,b,c,d`.

Emphasised observation. Both configurations solve all seven cells. Every paired `delta.aba` is distinct. The emphasised pattern is on the `sechk` side: in all seven `sechk` deltas, the only observed BK predictor that appears is the *first* predictor in that target's BK order. Later predictors are absent. Matched `relto` deltas cite that first predictor and also later ones:

Cell	BK order	Predictors in <code>sechk</code> delta	Extra predictors in <code>relto</code> delta
H0 <i>a</i>	<i>b, c</i>	<i>b</i>	<i>c</i>
H0 <i>b</i>	<i>a, c</i>	<i>a</i>	<i>c</i>
H0 <i>c</i>	<i>a, b</i>	<i>a</i>	<i>b</i>
H3 <i>a</i>	<i>b, c, d</i>	<i>b</i>	<i>c, d</i>
H3 <i>b</i>	<i>a, c, d</i>	<i>a</i>	<i>c, d</i>
H3 <i>c</i>	<i>a, b, d</i>	<i>a</i>	<i>b, d</i>
H3 <i>d</i>	<i>a, b, c</i>	<i>a</i>	<i>b, c</i>

For target `a` the first BK predictor is `b`. Illustrative H0 child `c` (BK order `a, b`):

```
% ecai2024 / relto
c(A) :- alpha_1(A), a_val_1(A).
c_alpha_1(A) :- b_val_0(A).

% nd_brave_sechk / sechk
c(A) :- alpha_1(A), a_val_1(A).
c_alpha_1(A) :- alpha_2(A), a_val_1(A).
c_alpha_2(A) :- alpha_1(A), a_val_1(A).
```

Both share `c(A) :- alpha_1(A), a_val_1(A)`. Under `relto`, the contrary of α_1 is `b_val_0`. Under `sechk`, every BK literal in the delta is an *a*-value literal.

Trace-supported explanation. After a fold fails post-folding entailment, the policies fork in `gen.pl`:

```
relto: gen2: ... looking for assumption relative to
sechk: gen2: generating NEW assumption: alpha_1(A)
```

Under `relto`, a residual reuse that returns `K0`, cannot introduce an assumption for ... can hand control back to folding, so a later BK literal (for example `b_val_0` on `H0 c`) may be tried next. That is how later predictors enter the `relto` deltas.

Under `sechk`, a provisional assumption is minted immediately and search continues through `gen5/gen6`. On these brave cells, the accepted theory is completed using only value literals of the first BK predictor, before those later folds are reached. Hence every `H7a sechk` delta cites exactly one observed BK predictor: the first in that target's BK order.

Evidential boundary. `H7a` is a single-option `asm_intro` ablation under brave `nd` on these seven pairs. The cautious counterpart is `H7b`.

Primary evidence record: `docs/experiments/qualitative/M13-C3-binary-collider-and/h7a_relto_vs_sechk.md`

5.8.2 H7b: cautious nd

Question and purpose. Under cautious `nd` on the locked `H0` and `H3` $n = 30$ tables, if only `asm_intro` changes from `relto` to `sechk`, which BK predictors appear in the accepted deltas for child `c`?

Controlled comparison. Fixture, frozen seed-42 sample, target `c`, examples, BK, `learning_mode(cautious)`, `nd` folding, `folding_selection(any)`, `folding_space(all)`, `folding_steps(10)`, `post_folding_test_entailm` and `check_ic` were held fixed. Locked `baseline_cautious` uses `asm_intro(relto)`. Experimental `nd_cautious_sechk` matches it except for `asm_intro(sechk)`. The lead pairs are `H0 c` and `H3 c`.

Emphasised observation. Both configurations solve both child cells. Every paired `delta.aba` is distinct. On both tables, the cautious `sechk` delta cites the first BK predictor and also later predictors. That is the emphasised contrast with `H7a`, where every brave `sechk` delta cited only the first BK predictor.

For `H0 c` (BK order a, b):

```
% baseline_cautious / relto
c(A) :- alpha_1(A), a_val_1(A).
c_alpha_1(A) :- b_val_0(A).

% nd_cautious_sechk / sechk
c(A) :- alpha_1(A), a_val_1(A).
c_alpha_1(A) :- alpha_2(A), a_val_1(A).
c_alpha_2(A) :- b_val_1(A).
```

Both share `c(A) :- alpha_1(A), a_val_1(A)`. Under `relto`, the contrary of α_1 is `b_val_0`. Under `sechk`, the contrary of α_1 contains a further assumption α_2 conditioned on `a_val_1`, and the contrary of α_2 is `b_val_1`. Parent b therefore appears in both deltas. The closing b -literals differ (`b_val_0` versus `b_val_1`), and `sechk` inserts one extra assumption between the shared target rule and that b -literal.

For H3 *c* (BK order *a, b, d*), both deltas keep ordinary target rules conditioned on leading *a* and place parents *b* and *d* in contrary rules. The **sechk** delta has five assumptions against three under **relto**, and its trace is longer (899 lines against 257).

Trace-supported explanation. The policies fork as in H7a after a fold fails post-folding entailment:

```
relto:  gen2: ... looking for assumption relative to
sechk:  gen2: generating NEW assumption: alpha_1(A)
```

Under cautious learning, residual reuse often returns KO, **cannot introduce an assumption** for ... in **gen3**. On the **relto** path, that KO returns control to folding, so a later BK literal such as **b_val_0** on H0 *c* can be tried next. On the **sechk** path, a provisional assumption is installed first and search continues through **gen5/gen6**. When an early candidate that uses only *a*-literals fails cautious checking, further assumption introduction and backtracking still reach later *b* or *d* predicates. That is why these cautious **sechk** deltas for *c* include mechanism parents, while the brave H7a **sechk** deltas for *c* used only the first BK predictor.

The same mint-first path accounts for the longer H3-*c* **sechk** trace. More provisional repairs are attempted before later folds succeed.

Evidential boundary. H7b is a single-option **asm_intro** ablation under cautious nd on the locked H0 and H3 child-*c* cells, with **baseline_cautious** as the **relto** comparator.

Primary evidence record: docs/experiments/qualitative/M13-C3-binary-collider-and/h7b-cautious_relto_vs_sechk.md

6 Findings

6.1 Finding 1: Greedy solutions retain all represented BK predictors

6.2 Finding 2: nd solutions are organised around the first BK predictor

6.3 Finding 3: brave nd can solve predictor-ambiguous root targets

6.4 Finding 4: the folding-token ceiling repeats one search attempt

Statement. Under nd folding, **folding_steps**(*M*) bounds how many times the learner restarts the whole learning problem. Search begins at **tokens**(1). When the attempt at token allowance *T* exhausts without an accepted theory and $T + 1 \leq M$, the learner raises the allowance to $T + 1$ and restarts on the same rules and the same $\langle E^+, E^- \rangle$. H6a calls one such attempt a failed band.

Under the exact-value target-wise encoding of Section 2.4, every folding call completes after one folding step. An allowance of one is therefore sufficient at every point of the search, the set of folded rules reachable at allowance $T + 1$ coincides with the set reachable at allowance *T*, and each restart repeats its predecessor. Two consequences hold for every target under this encoding.

1. A target that receives an accepted theory receives it inside the first allowance. Raising **folding_steps** leaves the emitted delta, the trace, and the runtime unchanged.
2. A target that receives no accepted theory produces *M* copies of one search attempt. Trace length and runtime are affine in *M*, and the outcome is **completed_no_solution** at every *M*.

The property that limits each folding call to one step belongs to the encoding, so the statement holds for both learning modes, for both `asm_intro` options, and for all fixtures encoded this way.

Supporting evidence. H6a measured the second consequence on the H0 roots under `baseline_cautious` at $M \in \{1, 2, 5, 10\}$, where $\text{lines}(M) = 51 + 498M$ and $N_{\text{asm}}(M) = 9M$ hold exactly at all four ceilings, and where control `c` is invariant in M .

Generality rests on an audit of the Prolog traces already recorded for this investigation, performed by `docs/experiments/qualitative/M13-C3-binary-collider-and/band_replay_audit.py`. No new learner runs were required. The audit covers every target-wise trace whose effective options include `folding_mode(nd)`. That is 49 cells: the 45 cells of the two fixtures described here, together with four earlier diamond-fixture runs retained as infrastructure provenance and used only as evidence about engine behaviour. The 49 cells span the four `nd` configurations of Section 4.1 and the three H6a `folding_steps` variants, brave and cautious learning, `relto` and `sechk`, and $n \in \{30, 50, 60, 90\}$.

Audited property of the recorded <code>nd</code> traces	Count
Target-wise traces with effective <code>folding_mode(nd)</code>	49
Atom lists offered to a folding call containing more than one atom	0
Folding calls that spent more than one folding token	0
Traces printing the token-exhaustion message	0
<code>solved</code> cells that raised the token allowance	0 of 30
<code>completed_no_solution</code> cells that raised the allowance	14 of 19
Cells whose search attempts are all identical after normalisation	14 of 14

The five `completed_no_solution` cells that raised no allowance are the `folding_steps(1)` configurations, for which M permits no restart. Attempt identity was checked by dividing each trace at its `* Increasing folding tokens to:` lines and comparing the resulting attempts after normalising the three quantities that legitimately change between them: generated assumption indices, internal Prolog variable numbers, and the printed token counter. For `baseline_cautious` target `a` at $M = 10$, the trace of 5031 lines divides into a 51-line preamble, ten attempts of 497 lines each, the nine restart announcements, and the closing `* No solution found!` line. All ten attempts are identical under that normalisation. Target `b` gives ten identical attempts of 506 lines.

The restart count multiplies whatever one attempt costs. Under `nd_cautious_sechk`, which differs from `baseline_cautious` only in `asm_intro(sechk)`, the H0 roots also perform ten identical attempts, of 1367 lines each for target `a` and 1376 lines each for target `b`, with runtimes of 167.6s and 172.9s. H7b records the same sequence of restarts from `tokens(2)` to `tokens(10)` on those cells and reports the resulting cost divergence. Across all 14 cells that raised the allowance, the size of one attempt ranges from 497 to 9747 lines, and attempt identity holds in every one of them.

For the first consequence, every accepted theory recorded in this investigation was accepted inside the first token allowance. All 30 solved `nd` cells therefore carry the deltas reported in H0, H3, H4, H4b, H7a, and H7b at `folding_steps(1)` as well.

Trace-supported explanation. The folding-token counter bounds folding steps inside a single folding call. One step replaces one atom of the body being folded by the head of a matching rule, and the call completes as soon as the list of atoms to be folded is empty.

Under the exact-value encoding both sides of that step carry a single atom. The rules offered to folding are rote rules over row identifiers, written in the trace as `ert: a(A) <- [A=1]`, and each background rule has the form `b_val_1(A) :- A=1`. Folding the single row-identifier equality of a

rote rule against such a background rule leaves nothing further to fold, so the call completes after one step. Every recorded nd trace shows the corresponding pair of lines, here at allowance $T = 2$:

```
2: folding [A=1] with b_val_1(_7934): b_val_1(A) <- [A=1]
1: DONE
```

The step is printed at the current allowance and completion is printed one below it, leaving $T - 1$ tokens unspent. The engine reaches its token-exhaustion diagnostic only when the counter is zero while atoms remain to be folded, and that combination arises in no recorded trace.

Because the reachable folded rules are the same at every allowance, a restarted search repeats its predecessor. The counter that names new assumptions is global and continues across restarts, so the same assumption-introduction events reappear under higher indices. That is the source of $N_{\text{asm}}(M) = 9M$ in H6a.

Evidential boundary. The condition under which the ceiling repeats a search is the exact-value encoding. Background knowledge containing a rule whose body has more than one atom would leave atoms to be folded after the first step, and a higher allowance would then admit folded bodies that a lower allowance cannot reach. The finding therefore concerns `folding_steps` together with this encoding.

The audit establishes the one-step property for the runs performed here. A proof about the folding predicate for arbitrary background knowledge was not attempted.

Trace length is exactly affine in M on the H0 roots. Runtime is affine to within measurement variation: per-attempt runtime for target `a` lies between 6.56 and 7.00s, giving an $M = 10$ to $M = 1$ runtime ratio of 10.54.

The root outcome `completed_no_solution` under cautious learning records the absence of a cautiously acceptable covering theory, as established in H4. Finding 4 concerns the cost of reaching that outcome.

Primary evidence record: docs/experiments/qualitative/M13-C3-binary-collider-and/h6_folding_and_n_ablation.md

6.5 Finding 5: cautious-nd trace size grows with grounded sample size

6.6 Finding 6: relto and sechk produce different assumption-introduction paths

7 Implications